

Homework 2: Probability Models

BEE 4850/5850

Name: Luca Katzen

ID: 5360358

Due Date

Friday, 2/20/26, 9:00pm

Overview

Instructions

The goal of this homework assignment is to practice developing and working with probability models for data.

- Problem 1 asks you to fit some linear regressions to Chicago mortality data.
- Problem 2 asks you to analyze the probability of Cayuga Lake freezing using a logistic regression model.
- Problem 3 asks you to use Poisson regression to predict salamander counts based on environmental data.
- Problem 4 asks you to look at the impact of the gender of hurricane names on deaths^[1].

[1] Yes, seriously. Ish. Trust me, I know.

Load Environment

The following code loads the environment and makes sure all needed packages are installed. This should be at the start of most Julia scripts.

```
import Pkg
Pkg.activate(@__DIR__)
Pkg.instantiate()
```

Activating project at `~/Documents/GitHub/hw02-LucaKatzen1-Spring26`

Installed GR_jll	v0.73.22+0
Installed Opus_jll	v1.6.1+0
Installed JpegTurbo_jll	v3.1.4+0
Installed GR	v0.73.22
Installed PDMats	v0.11.37
Installed Preferences	v1.5.1
Installed Libmount_jll	v2.41.3+0
Installed StaticArrays	v1.9.16
Installed EnumX	v1.0.6
Installed FiniteDiff	v2.29.0
Installed ADTypes	v1.21.0
Installed MuladdMacro	v0.2.4
Installed JSON	v1.4.0
Installed XZ_jll	v5.8.2+0
Installed NLSolversBase	v8.0.0
Installed StatsBase	v0.34.10
Installed NearestNeighbors	v0.4.27
Installed libpng_jll	v1.6.54+0
Installed PositiveFactorizations	v0.2.4
Installed Conda	v1.10.3
Installed Plots	v1.41.5
Installed IJulia	v1.34.1
Installed ZMQ	v1.5.1
Installed KernelDensity	v0.6.11
Installed ArrayInterface	v7.22.0
Installed OpenSSL_jll	v3.5.5+0
Installed LineSearches	v7.6.0
Installed FFMPEG_jll	v8.0.1+0
Installed DifferentiationInterface	v0.7.16
Installed Libuuid_jll	v2.41.3+0
Installed IntegerMathUtils	v0.1.3
Installed Optim	v2.0.1
Installed FillArrays	v1.16.0
Installed StructUtils	v2.6.2
Installed Primes	v0.5.7
Installed FFTA	v0.3.1
Installed Glib_jll	v2.86.3+0

```

Installed WoodburyMatrices      v1.1.0
Installed Distributions          v0.25.123
Building Conda → `~/.julia/scratchspaces/44cfe95a-1eb2-52ea-b672-e2afdf69b78f/8f06b0cfa4
Building IJulia → `~/.julia/scratchspaces/44cfe95a-1eb2-52ea-b672-e2afdf69b78f/df8af3ab
Precompiling project...
381.3 ms  MuladdMacro
307.6 ms  EnumX
308.5 ms  IntegerMathUtils
400.5 ms  PositiveFactorizations
462.9 ms  ArrayInterface
522.1 ms  ADTypes
684.9 ms  WoodburyMatrices
732.7 ms  FillArrays
818.6 ms  PDMats
552.8 ms  TableOperations
559.4 ms  Preferences
631.1 ms  StructUtils
369.2 ms  ArrayInterface → ArrayInterfaceStaticArraysCoreExt
669.3 ms  Primes
499.7 ms  ArrayInterface → ArrayInterfaceChainRulesCoreExt
732.2 ms  ArrayInterface → ArrayInterfaceSparseArraysExt
501.9 ms  ADTypes → ADTypesConstructionBaseExt
508.7 ms  ADTypes → ADTypesChainRulesCoreExt
622.0 ms  AxisAlgorithms
582.0 ms  FillArrays → FillArraysStatisticsExt
559.3 ms  FillArrays → FillArraysSparseArraysExt
473.3 ms  PrecompileTools
1082.0 ms  DifferentiationInterface
641.0 ms  JLLWrappers
690.2 ms  StructUtils → StructUtilsTablesExt
512.7 ms  FiniteDiff
895.1 ms  FillArrays → FillArraysPDMatsExt
1777.7 ms  StatsBase
697.0 ms  DifferentiationInterface → DifferentiationInterfaceChainRulesCoreExt
861.2 ms  DifferentiationInterface → DifferentiationInterfaceSparseArraysExt
890.5 ms  OpenSSL_jll
1255.6 ms  RecipesBase
1445.5 ms  FFTA
1340.2 ms  StringManipulation
535.9 ms  Graphite2_jll
419.9 ms  Libmount_jll
398.4 ms  EpollShim_jll
417.9 ms  Bzip2_jll

```

451.4 ms	LLVMOpenMP_jll
424.0 ms	libsodium_jll
431.8 ms	Rmath_jll
377.4 ms	Xorg_libICE_jll
388.8 ms	Xorg_libXau_jll
403.6 ms	libfdk_aac_jll
442.9 ms	libpng_jll
464.5 ms	LAME_jll
427.5 ms	LERC_jll
452.4 ms	fzf_jll
540.5 ms	JpegTurbo_jll
441.2 ms	Ogg_jll
491.8 ms	XZ_jll
416.0 ms	mtdev_jll
397.5 ms	Xorg_libXdmcp_jll
423.2 ms	x265_jll
424.3 ms	x264_jll
487.4 ms	libaom_jll
446.7 ms	Zstd_jll
552.5 ms	Arpack_jll
497.6 ms	Expat_jll
470.2 ms	LZO_jll
443.6 ms	Opus_jll
419.6 ms	Xorg_xtrans_jll
447.8 ms	libevdev_jll
384.8 ms	Libffi_jll
3404.7 ms	ColorSchemes
479.9 ms	Libiconv_jll
452.1 ms	eudev_jll
443.0 ms	OpenSpecFun_jll
414.3 ms	Libuuid_jll
404.1 ms	FriBidi_jll
517.1 ms	FiniteDiff → FiniteDiffSparseArraysExt
581.0 ms	DifferentiationInterface → DifferentiationInterfaceFiniteDiffExt
753.3 ms	PDMats → StatsBaseExt
350.7 ms	Xorg_libSM_jll
1002.2 ms	FreeType2_jll
4964.1 ms	StaticArrays
1097.8 ms	Pixman_jll
1116.8 ms	ZeroMQ_jll
350.3 ms	Xorg_libxcb_jll
1301.7 ms	Rmath
437.4 ms	DBus_jll

1285.9 ms	JLFzf
1197.6 ms	libvorbis_jll
1030.9 ms	Arpack
537.3 ms	Wayland_jll
2814.8 ms	OpenSSL
383.6 ms	libinput_jll
789.4 ms	GettextRuntime_jll
2018.6 ms	Ghostscript_jll
947.2 ms	NLSolversBase
7514.6 ms	Parsers
528.7 ms	Adapt → AdaptStaticArraysExt
630.0 ms	ConstructionBase → ConstructionBaseStaticArraysExt
2870.8 ms	Libtiff_jll
1136.3 ms	StaticArrays → StaticArraysStatisticsExt
1154.6 ms	StaticArrays → StaticArraysChainRulesCoreExt
1673.3 ms	Fontconfig_jll
466.3 ms	Xorg_xcb_util_jll
560.4 ms	DifferentiationInterface → DifferentiationInterfaceStaticArraysExt
559.2 ms	FiniteDiff → FiniteDiffStaticArraysExt
572.7 ms	FillArrays → FillArraysStaticArraysExt
2220.0 ms	SpecialFunctions
392.3 ms	Xorg_libX11_jll
677.5 ms	LineSearches
841.5 ms	InlineStrings → ParsersExt
1980.2 ms	ZMQ
1539.4 ms	Glib_jll
430.1 ms	Xorg_xcb_util_image_jll
494.6 ms	Xorg_xcb_util_keysyms_jll
544.7 ms	Xorg_xcb_util_renderutil_jll
656.7 ms	Xorg_xcb_util_wm_jll
2756.2 ms	JSON
3289.1 ms	Latexify
2563.9 ms	Interpolations
423.5 ms	Xorg_libXrender_jll
1032.4 ms	HypergeometricFunctions
723.7 ms	ColorVectorSpace → SpecialFunctionsExt
506.8 ms	Xorg_libXext_jll
478.6 ms	Xorg_libXfixes_jll
1836.2 ms	SpecialFunctions → SpecialFunctionsChainRulesCoreExt
522.8 ms	Xorg_libxkbfile_jll
6885.0 ms	PlotUtils
506.7 ms	Xorg_xcb_util_cursor_jll
4158.8 ms	NearestNeighbors

327.0 ms	Xorg_libXinerama_jll
1079.1 ms	WeakRefStrings
926.2 ms	Conda
340.4 ms	Xorg_libXrandr_jll
1097.6 ms	Latexify → SparseArraysExt
375.3 ms	Libglvnd_jll
378.5 ms	Xorg_libXi_jll
379.9 ms	Xorg_libXcursor_jll
371.1 ms	Xorg_xkbcomp_jll
1994.0 ms	StatsFuns
2637.5 ms	Clustering
2832.6 ms	PlotThemes
822.5 ms	Xorg_xkeyboard_config_jll
3488.1 ms	RecipesPipeline
4031.0 ms	Cairo_jll
532.9 ms	xkbcommon_jll
539.2 ms	Vulkan_Loader_jll
6390.9 ms	Optim
1976.1 ms	StatsFuns → StatsFunsChainRulesCoreExt
1149.1 ms	HarfBuzz_jll
1199.5 ms	libass_jll
1223.6 ms	Pango_jll
502.0 ms	libdecor_jll
12268.5 ms	HTTP
439.4 ms	GLFW_jll
3398.4 ms	Qt6Base_jll
4881.2 ms	Distributions
3054.0 ms	FFMPEG_jll
1272.2 ms	Qt6ShaderTools_jll
1173.4 ms	Distributions → DistributionsChainRulesCoreExt
1232.4 ms	Distributions → DistributionsTestExt
1692.5 ms	MultivariateStats
1909.0 ms	Qt6Declarative_jll
2608.7 ms	FFMPEG
369.3 ms	Qt6Wayland_jll
10884.8 ms	IJulia
22164.7 ms	PrettyTables
11959.1 ms	CSV
3238.0 ms	GR_jll
2969.8 ms	KernelDensity
2261.4 ms	GR
1447.8 ms	GR → IJuliaExt
23440.1 ms	DataFrames

```
1443.0 ms    Latexify → DataFramesExt
33088.2 ms   Plots
2561.3 ms    Plots → IJuliaExt
5041.9 ms    StatsPlots
171 dependencies successfully precompiled in 70 seconds. 122 already precompiled.
```

The following packages are included in the environment (to help you find other similar packages in other languages). The code below loads these packages for use in the subsequent notebook (the desired functionality for each package is commented next to the package).

```
using Random # random number generation and seed-setting
using DataFrames # tabular data structure
using CSV # reads/writes .csv files
using Distributions # interface to work with probability distributions
using Plots # plotting library
using StatsBase # statistical quantities like mean, median, etc
using StatsPlots # some additional statistical plotting tools
using Optim # optimization tools
using LaTeXStrings # latex formatting for plot strings
```

Problems

Problem 1 (6 points)

Let's revisit the `chicago` dataset from HW1 (found in `data/chicago.csv`). We will look at the relationship of the potential predictor variables `pm25median` (the median density of smaller pollutant particles), `o3median` (the median concentration of O_3), `so2median` (the median concentration of SO_2), and `tmpd` (mean daily temperature) with the variable we would like to predict, `death` (the number of non-accidental deaths on that day).

Problem 1.1

Plot each predictor variable against the response variable (make sure to include labels and titles for each plot). Describe what you observe about the bivariate relationships, if one exists (*e.g.* does it look like they exist, are they linear, positive/negative, etc.) Does a linear regression seem appropriate from any of these plots? Which predictor do you think is most appropriate for a linear regression model and why?

Problem 1.2

You decide to focus on the relationship between `tmpd` and `death`. Your classmate suggests that you use a Gaussian-error linear regression model. Based on this suggestion, calculate the regression coefficients and error variance. Explain the interpretation of the coefficients that you obtain **in this specific problem context**.

Problem 1.3

Add your fitted regression line and a 90% prediction interval to your scatterplot of `tmpd` and `death`. Using this and any other diagnostics, do you agree with the Gaussian-error linear model assumption? Why or why not?

```
#Problem 1.1
data = DataFrame(CSV.File("data/chicago.csv"))

p1 = scatter(
  data.pm25median,
  data.death,
  xlabel = "PM25 Median",
  ylabel = "Number of Deaths",
  title = "PM25 Median vs. Deaths",
  legend = false
)

p2 = scatter(
  data.o3median,
  data.death,
  xlabel = "o3 Median",
  ylabel = "Number of Deaths",
  title = "o3 Median vs. Deaths",
  legend = false
)

p3 = scatter(
  data.so2median,
  data.death,
  xlabel = "so2 Median",
  ylabel = "Number of Deaths",
  title = "so2 Median vs. Deaths",
  legend = false
)
```




```

p4 = scatter(
    data.tmpd,
    data.death,
    xlabel = "Mean Daily Temperature",
    ylabel = "Number of Deaths",
    title = "Temperature vs. Deaths",
    legend = false
)

plot(p1, p2, p3, p4, layout = (2, 2))

```



Problem 1.1

None of the graphs appear to have strong or obvious bivariate correlations. While O3 displays some outliers at higher quantities, PM25, O3, and SO2 don't exhibit distinct linear patterns with deaths (the dots are dispersed without much visible direction). The most promising predictor is temperature, which exhibits a significant correlation with increases in mortality, particularly at the upper end of the temperature range. Temperature is the best candidate for a basic linear regression model since it exhibits a directional pattern that is worth modeling,

even though there is still a lot of scatter. It also has the most pronounced linear trend of the four variables.

```
#Problem 1.2
X = hcat(ones(length(data.tmpd)), data.tmpd)
y = data.death

Beta = (X' * X) \ (X' * y)

yline = X * Beta
residuals = y - yline

n = length(y)
sigmasquared = sum(residuals.^2) / (n - 2)

println("Intercept: ", Beta[1])
println("Slope: ", Beta[2])
println("Error variance: ", sigmasquared)
```

```
Intercept: 129.95705121476817
Slope: -0.2896442961500982
Error variance: 202.31531425938184
```

Problem 1.2

The intercept is the expected number of deaths on a day when the mean daily temperature is 0°F. The slope is the expected change in the number of deaths for each 1°F increase in mean daily temperature. Because β_2 is negative, warmer days are not associated with more deaths. The error variance (σ^2) represents how much the daily death counts vary around a linear trend, capturing variability due to other factors besides temperature. The variance is roughly 202 so predictions tend to be off by around 14 deaths.

```
#Problem 1.3
X = hcat(ones(length(data.tmpd)), data.tmpd)
y = data.death
beta = (X' * X) \ (X' * y)
yline = X * beta
residuals = y - yline
n = length(y)
sigmasquared = sum(residuals.^2) / (n - 2)
sigma = sqrt(sigmasquared)
```


```

tmpdrange = range(minimum(data.tmpd), maximum(data.tmpd), length=100)
xpred = hcat(ones(length(tmpdrange)), tmpdrange)
ypred = xpred * beta

tval = quantile(TDist(n-2), 0.95)
sepred = sigma .* sqrt.(1 .+ sum((xpred * inv(X' * X)) .* xpred, dims=2))
predlower = ypred .- tval .* vec(sepred)
predupper = ypred .+ tval .* vec(sepred)

scatter(
    data.tmpd,
    data.death,
    xlabel = "Mean Daily Temperature",
    ylabel = "Number of Deaths",
    title = "Temperature vs. Deaths",
    label = "Data",
    alpha = 0.5
)
plot!(tmpdrange, ypred,
    label = "Fitted Line",
    color = :red,
    linewidth = 2)
plot!(tmpdrange, predlower,
    label = "90% Prediction Interval",
    color = :red,
    linestyle = :dash)
plot!(tmpdrange, predupper,
    label = "",
    color = :red,
    linestyle = :dash)

```



Problem 1.3

I don't fully agree with the Gaussian-error linear model assumption. There are clear outliers with deaths exceeding 200-400 that fall well outside the 90% prediction interval, especially at higher temperatures on the scatter plot. The relationship also appears to have variability in deaths increases at higher temperatures. While overall there's a negative linear trend, the extreme outliers and non-constant variance suggest that the Gaussian assumption is not ideal.

Problem 2 (6 points)

This problem is adapted from Example 4.1 and 4.2 in Daniel Wilks' *Statistical Methods in the Atmospheric Sciences*. Updating the table in that book, Cayuga Lake has frozen in 1796, 1816, 1856, 1875, 1884, 1904, 1912, 1934, 1961, 1979, and 2015[1]. The probability of whether a severe enough cold spell to freeze the lake occurs is related to changes in global mean temperature. The file `data/HadCRUT5.1Analysis_gl.txt` contains estimates of monthly global mean temperature anomalies (in °C, relative to the 1961–1990 mean) as well as the annual global temperature anomaly. You would like to build a model to predict the probability that Cayuga Lake freezes based on the winter (DJF) temperature anomaly.

Problem 2.1

Write down a logistic regression model for the probability that Cayuga Lake freezes in a given winter. Encode the occurrence of freezing as a 1, and non-freezing as 0. Load the temperature data and fit this model. The format of the data file is a little awkward, as the even rows are the number of stations (which should be ignored) and have uneven number of space delimiters. The first 13 columns are the year and the monthly means; the final column is the annual average (which we don't want to use). Another complication is that the (even-numbered) station rows don't have an entry in the 14th column. In Julia, you can load the file with

```
temp_dat = CSV.read("data/HadCRUT5.1Analysis_gl.txt", delim=" ",
    ignorerepeated=true, header=false,
    silencewarnings=true, DataFrame)
```

To simplify the analysis, since we don't always know whether the lake froze in the January starting a year or the December ending a year, let's assume that the years in which the lake froze correspond to the winter starting the year (so the DJF mean corresponding to the freezing in 1884 would include the December from 1883 and the January and February from 1884). You will only be able to use freezing events between 1851 and the present day due to the temperature data. Interpret the coefficients from your model fit **in this specific problem context**.

Problem 2.2

Plot the modeled probability of freezing as a function of the DJF temperature anomaly, and how this probability has changed over time. Does this model seem reasonable? How might you determine this?

Problem 2.3

What winter temperature anomaly would be required for less than a 1% probability of Cayuga Lake freezing?

```
data2 = CSV.read("data/HadCRUT5.1Analysis_gl.txt", delim=" ",
    ignorerepeated=true, header=false,
    silencewarnings=true, DataFrame)

oddrows = data2[1:2:end, :]
years = oddrows[2:end, 1]

djf = zeros(length(years))
for i in 1:length(years)
```

```

    dec_prev = oddrows[i, 13]
    jan_curr = oddrows[i+1, 2]
    feb_curr = oddrows[i+1, 3]
    djf[i] = (dec_prev + jan_curr + feb_curr) / 3
end

freezeyears = [1796, 1816, 1856, 1875, 1884, 1904, 1912, 1934, 1961, 1979, 2015]
freezeyears = freezeyears[freezeyears .>= 1851]

frozen = [year in freezeyears ? 1 : 0 for year in years]

logregdata = DataFrame(year = years, djf = djf, frozen = frozen)

println("DJF values when frozen: ", djf[frozen .== 1])

function logisticregression(X, y)
    function negloglikelihood(beta)
        eta = X * beta
        p = 1 ./ (1 .+ exp.(-eta))
        p = clamp.(p, 1e-10, 1 - 1e-10)
        return -sum(y .* log.(p) .+ (1 .- y) .* log.(1 .- p))
    end

    result = optimize(negloglikelihood, [0.0, -1.0], BFGS())
    return Optim.minimizer(result)
end

X = hcat(ones(length(logregdata.djf)), logregdata.djf)
y = logregdata.frozen

beta = logisticregression(X, y)

println("Intercept: ", beta[1])
println("Slope: ", beta[2])

```

DJF values when frozen: [-0.2793333333333333, -0.445, -0.37966666666666665, -0.7573333333333333]
 Intercept: -3.0163797707713482
 Slope: -0.714206431276419

Problem 2.1

is the log-odds of Cayuga Lake freezing when the DJF temperature anomaly is 0°C (at the 1961-1990 baseline). is how the log-odds of freezing changes per 1°C increase in the DJF temperature anomaly. Since is negative, this means warmer winter temperatures decrease the probability that the lake will freeze. Because is roughly equal to -0.7, each 1°C increase in temperature anomaly reduces the log-odds of freezing by 0.7, or a moderate decrease in freezing probability.


```
#Problem 2.2
djfrange = range(minimum(logregdata.djf), maximum(logregdata.djf), length=100)
xpred = hcat(ones(length(djfrange)), djfrange)
logodds = xpred * beta
probfreeze = 1 ./ (1 .+ exp.(-logodds))

p1 = scatter(logregdata.djf, logregdata.frozen,
    xlabel = "DJF Temperature Anomaly (°C)",
    ylabel = "Probability of Freezing",
    title = "Freezing Probability vs. Temperature",
    alpha = 0.5,
    label = "Observed")
plot!(djfrange, probfreeze,
    linewidth = 2,
    label = "Model",
    ylim = (0, 1))

xpredtime = hcat(ones(length(logregdata.djf)), logregdata.djf)
logoddstime = xpredtime * beta
probfreetime = 1 ./ (1 .+ exp.(-logoddstime))

p2 = scatter(logregdata.year, logregdata.frozen,
    xlabel = "Year",
    ylabel = "Probability of Freezing",
    title = "Freezing Probability Over Time",
    alpha = 0.5,
    label = "Observed")
plot!(logregdata.year, probfreetime,
    linewidth = 2,
    label = "Model",
    ylim = (0, 1))

plot(p1, p2, layout = (1, 2), size = (1000, 400))
```



Problem 2.2

The model seems reasonable overall. The probability of freezing decreases as temperature anomalies increase, which matches what you would expect logically. The probability of freezing over time shows a decline, reflecting the understood impacts of global warming. However, the model might be oversimplified since it only uses one predictor and there are relatively few freezing events to train it on. To assess the model we could check how well it classifies historical freezing events, calculate metrics like accuracy, or use cross-validation to test its predictive performance.

```
function findtemperature(beta, targetprob)
    return (log(targetprob / (1 - targetprob)) - beta[1]) / beta[2]
end

temp1 = findtemperature(beta, 0.01)
println("DJF temperature anomaly for 1% freezing probability: ", temp1, "°C")
```

DJF temperature anomaly for 1% freezing probability: 2.2104814661802203°C

Problem 2.3

For less than a 1% probability of Cayuga Lake freezing, the DJF temperature anomaly would need to exceed approximately 2.2°C.

Problem 3 (6 points) The file `data/salamanders.csv` contains counts of salamanders from 47, different plots of the same area in California, as well as the percentage of ground cover and age of the forest in the plot. You would like to see if you can use these data to predict the salamander counts with a Poisson regression.

Problem 3.1

Write down and fit a Poisson regression model for salamander counts using the percentage of ground cover as a predictor. You may need to standardize the predictors as they are much larger than the counts, which you can do using the following function (or its equivalent): [1] Many of these dates, particularly prior to the 20th century, are subject to some debate, and would be different depending on the data source. Nevertheless, let's use this dataset.

Problem 3.2

Plot the expected counts and 90% prediction intervals from your model (based on 1,000 simulations from the model) along with the data. How well does the model predict the observed counts? In what ways does it do a good or bad job?

Problem 3.3

Can you improve the model by including forest age as a predictor? How did you determine whether this helps or does not help with prediction?

```
#Problem 3.1
data3 = DataFrame(CSV.File("data/salamanders.csv"))

function standardize(x)
    return (x .- mean(x)) ./ std(x)
end

data3.pct_cover_std = standardize(data3.PCTCOVER)

function poissonregression(X, y)
    function negloglikelihood(beta)
        lambda = exp.(X * beta)
        return -sum(y .* log.(lambda) .- lambda)
    end

    result = optimize(negloglikelihood, zeros(size(X, 2)), BFGS())
    return Optim.minimizer(result)
end

X = hcat(ones(nrow(data3)), data3.pct_cover_std)
y = data3.SALAMAN

beta = poissonregression(X, y)
```

```
println("Intercept ( ): ", beta[1])
println("Slope ( ): ", beta[2])
```

```
Intercept ( ): 0.42951129115472364
Slope ( ): 1.159506131879513
```

Problem 3.1

is the log expected salamander count when percent cover is at its mean value. is the change in log expected count for each one standard deviation increase in percent cover. Because is positive, more ground cover is associated with higher salamander counts. This makes sense because salamanders need moist habitat and more ground cover means more shade and moisture.

```
#Problem 3.2
coverrange = range(minimum(data3.pct_cover_std), maximum(data3.pct_cover_std), length=100)
xpred = hcat(ones(length(coverrange)), coverrange)
lambda = exp.(xpred * beta)

nsim = 1000
sims = zeros(length(coverrange), nsim)
for i in 1:length(coverrange)
    for j in 1:nsim
        sims[i, j] = rand(Poisson(lambda[i]))
    end
end


predlower = [quantile(sims[i, :], 0.05) for i in 1:length(coverrange)]
predupper = [quantile(sims[i, :], 0.95) for i in 1:length(coverrange)]

scatter(data3.pct_cover_std, data3.SALAMAN,
        xlabel = "Standardized Percent Cover",
        ylabel = "Salamander Count",
        title = "Salamander Population vs. Ground Cover",
        label = "Observed",
        alpha = 0.6)
plot!(coverrange, lambda,
        label = "Expected Count",
        linewidth = 2,
        color = :red)
plot!(coverrange, predlower,
```

```

label = "90% Prediction Interval",
color = :red,
linestyle = :dash)
plot!(coverrange, predupper,
label = "",
color = :red,
linestyle = :dash)

```



Problem 3.2

The model does an okay job at portraying the general trend of salamander counts with ground cover, but has some significant limitations. The most obvious issue is that many observed counts fall outside the 90% prediction intervals, especially at higher cover values. The biggest problem is overdispersion as the observed data is much more spread out than the Poisson model expects. Since Poisson assumes $\text{variance} = \text{mean}$, it underestimates the true variability in salamander counts. The wide scatter of points at a given cover level suggests other unmeasured factors are influencing counts, or that a different distribution would better model the variability in the data.

```

#Problem 3.3
data3.forest_age_std = standardize(data3.FORESTAGE)

X2 = hcat(ones(nrow(data3)), data3.pct_cover_std, data3.forest_age_std)
beta2 = poissonregression(X2, y)

println("Model with forest age:")
println("Intercept: ", beta2[1])
println("Slope for cover: ", beta2[2])
println("Slope for age: ", beta2[3])

lambda1 = exp(X * beta)
lambda2 = exp(X2 * beta2)

loglik1 = sum(y .* log.(lambda1) .- lambda1)
loglik2 = sum(y .* log.(lambda2) .- lambda2)

aic1 = -2 * loglik1 + 2 * length(beta)
aic2 = -2 * loglik2 + 2 * length(beta2)

println("\nModel 1 (cover only) AIC: ", aic1)
println("Model 2 (cover + age) AIC: ", aic2)
println("AIC improvement: ", aic1 - aic2)

```

```

Model with forest age:
Intercept: 0.4293187208268531
Slope for cover: 1.1624676135180043
Slope for age: -0.00408443026687972

Model 1 (cover only) AIC: -42.512375571960845
Model 2 (cover + age) AIC: -40.514172408183775
AIC improvement: -1.9982031637770703

```

Problem 3.3

I determined whether forest age can help to improve the model by comparing AIC values. Based on this, adding forest age does not appear to improve the model. The AIC actually increases slightly (gets worse) by about 2 when we add forest age as a predictor, suggesting the additional complexity isn't actually providing better fit. The slope for forest age is also very small (roughly -0.004), indicating that it has minimal effect on salamander counts. This suggests that once ground cover is accounted for, forest age doesn't provide additional useful information for predicting salamander counts, so the original model makes more sense.

Problem 4 (7 points)

In 2014, [a paper was published in a prestigious journal](#) which claimed that hurricanes with more feminine names are deadlier than hurricanes with more masculine names because people take warnings about female-named hurricanes less seriously[1]. The file `data/Hurricanes.csv` contains the original data used in this analysis (note that this dataset excludes Hurricanes Katrina and Audrey because they were “outliers”). While we won’t replicate the specific analysis in this paper, let’s use the data to look at this hypothesis.

Problem 4.1

One might interpret the hypothesis to claim that the impact of the name is strengthened by the more powerful. A measure of hurricane strength is its minimum pressure (`min_pressure` in the dataset). Fit a Poisson regression model that predicts deaths (`deaths`) using the gender of the name (`female`) and minimum pressure (you may need to standardize the pressure). We use `female` instead of the continuous measure of the `femininity` of the name as this is highly subjective and has some weird coding (such as “Sandy” being coded as one of the most female names).

Problem 4.2

Interpret the results by generating 10,000 counterfactual simulations for hurricanes with male and female names. Plot the expected values and 90% prediction intervals from these two sets of simulations and compare with the observed storm deaths. Where does the model do well or not well?

Problem 4.3

Does the gender effect size resulting from this model seem reasonable to you? Draw on qualitative or quantitative assessments to justify your assessment of reasonableness.

Problem 4.4

We can now stop taking this hypothesis seriously. To what extent do you think the finding could be an artifact of the dataset (e.g. there is no actual effect, but there are coincidental features of the data that produce the result that female-named storms are more deadly than male-named storms)? Justify this conclusion with specific reference to an exploratory analysis of the data.

References

[1] This paper has become a bit of a joke among statisticians, but let's take the hypothesis seriously for this problem's sake.

```
#Problem 4.1
data4 = DataFrame(CSV.File("data/hurricanes.csv"))

data4.min_pressure_std = standardize(data4.min_pressure)

X = hcat(ones(nrow(data4)), data4.female, data4.min_pressure_std)
y = data4.deaths

beta = poissonregression(X, y)

println("Intercept: ", beta[1])
println("Slope for female: ", beta[2])
println("Slope for pressure: ", beta[3])
```

```
Intercept: 2.4146672367810664
Slope for female: 0.4721117391418284
Slope for pressure: -0.719532469950491
```

Problem 4.1

The intercept is the log expected deaths for a male-named hurricane at average minimum pressure. The slope coefficient for female is the change in log expected deaths when comparing female to male-named hurricanes, holding pressure constant. Because the intercept is positive, female-named hurricanes are related to a higher fatality rate. The coefficient for pressure shows how log expected deaths change with each standard deviation change in minimum pressure. This coefficient is negative because lower pressure means stronger storms and more deaths.

```
nsim = 10000

pressurevalues = rand(data4.min_pressure_std, nsim)

Xmale = hcat(ones(nsim), zeros(nsim), pressurevalues)
lambdamale = exp.(Xmale * beta)
simsmale = [rand(Poisson(lambdamale[i])) for i in 1:nsim]

Xfemale = hcat(ones(nsim), ones(nsim), pressurevalues)
lambdafemale = exp.(Xfemale * beta)
```

```

simsfemale = [rand(Poisson(lambdafemale[i])) for i in 1:nsim]

expectedmale = mean(simsmale)
expectedfemale = mean(simsfemale)
pimale = quantile(simsmale, [0.05, 0.95])
pifemale = quantile(simsfemale, [0.05, 0.95])

println("Male-named hurricanes:")
println("Expected deaths: ", expectedmale)
println("90% Prediction Interval: ", pimale)

println("\nFemale-named hurricanes:")
println("Expected deaths: ", expectedfemale)
println("90% Prediction Interval: ", pifemale)

histogram(simsmale, alpha=0.5, label="Male", bins=0:1:maximum([simsmale; simsfemale]))
histogram!(simsfemale, alpha=0.5, label="Female", bins=0:1:maximum([simsmale; simsfemale]))
xlabel!("Deaths")
ylabel!("Frequency")
title!("Predicted Deaths by Hurricane Name Gender")

```

```


Male-named hurricanes:
Expected deaths: 14.6968
90% Prediction Interval: [3.0, 36.0]

```

```

Female-named hurricanes:
Expected deaths: 23.6239
90% Prediction Interval: [5.0, 56.0]

```



Problem 4.2

This graph shows that female-named hurricanes tend to have higher expected deaths than male-named ones. However, the prediction intervals are very wide and have lots of overlap, indicating that this relationship is not likely causal. The model struggles with extreme right-skewness of the data as most hurricanes cause few deaths while a handful cause catastrophic casualties. The Poisson distribution doesn't adequately capture this overdispersion and the rare extreme events, but it does a good job at predicting the low-fatality storms.

Problem 4.3

The gender effect size is not reasonable. The model predicts female-named hurricanes cause about 60% more deaths on average ($\exp(0.47)$ equals roughly 1.6), but on the histogram, the distributions are highly overlapping with male-named storms only shifted slightly to the right. Both distributions are right-skewed, concentrated near 0-20 deaths with long tails extending to 150+ deaths. The proposed hypothesis that implicit bias against female names causes people to ignore warnings would need to be backed by other data/variables to prove its relationship with storm casualties. It's highly unlikely that name gender alone would cause such a substantial

behavioral difference in people that they would override their self-preservation instincts during a dangerous hurricane.


```
#Problem 4.4
println("Male-named storms:")
println("Number: ", sum(data4.female .== 0))
println("Mean deaths: ", mean(data4[data4.female .== 0, :deaths]))
println("Median deaths: ", median(data4[data4.female .== 0, :deaths]))
println("Max deaths: ", maximum(data4[data4.female .== 0, :deaths]))

println("Female-named storms:")
println("Number: ", sum(data4.female .== 1))
println("Mean deaths: ", mean(data4[data4.female .== 1, :deaths]))
println("Median deaths: ", median(data4[data4.female .== 1, :deaths]))
println("Max deaths: ", maximum(data4[data4.female .== 1, :deaths]))

println("Year distribution:")
println("Male storms mean year: ", mean(data4[data4.female .== 0, :year]))
println("Female storms mean year: ", mean(data4[data4.female .== 1, :year]))

scatter(data4.year, data4.deaths,
        group=data4.female,
        xlabel="Year",
        ylabel="Deaths",
        title="Hurricane Deaths vs. Time",
        labels=["Male" "Female"],
        alpha=0.6)
```

```
Male-named storms:
Number: 30
Mean deaths: 14.233333333333333
Median deaths: 5.0
Max deaths: 84
Female-named storms:
Number: 62
Mean deaths: 23.758064516129032
Median deaths: 5.0
Max deaths: 256
Year distribution:
Male storms mean year: 1990.3333333333333
Female storms mean year: 1978.0967741935483
```



Problem 4.4

The finding is likely an artifact of the dataset. The most obvious issue visible in this plot is that all hurricanes prior to 1979 had female names, and those early storms were clearly far deadlier. This makes sense logically and historically as before 1979, warning systems, forecasting technology, etc. were dramatically worse, meaning storms caused more deaths regardless of names. Once alternating male/female naming began in 1979, the death tolls for both genders look much more comparable and generally lower. So the idea that female names cause more deaths is really just that older storms were deadlier. The paper's failure to account for this historical change in predictive technology and hurricane infrastructure is its fundamental flaw.